

Aurora AI API Reference

Base URL

`https://api.aurora.ai/v1`

All endpoints are served over HTTPS. Plain HTTP is not supported.

1. Authentication

Aurora uses **API keys** for authentication.

1.1 API Key Format

API keys are opaque strings, e.g.:

`aur_sk_live_xxxxxyyzzz...`

- Keys are **tenant-scoped** and can have:
 - `read`
 - `write`
 - `admin` scopes.

1.2 Authorization Header

Include the API key in the `Authorization` header:

`Authorization: Bearer aur_sk_live_xxxxxyyzzz`

1.3 Error Responses for Auth

- **401 Unauthorized** – missing or invalid key
- **403 Forbidden** – key does not have permission for requested action

Example error payload:

```
{  
  "error": {  
    "code": "unauthorized",  
    "message": "Invalid or missing API key.",  
    "details": null  
  }  
}
```

2. Models

2.1 List Models

GET /models

Returns metadata about available Aurora models.

Request

```
GET /v1/models HTTP/1.1  
Host: api.aurora.ai  
Authorization: Bearer aur_sk_live_xxx
```

Response

```
{  
  "data": [  
    {  
      "id": "aurora-72b-agent-v3",  
      "type": "text-completion",  
      "context_window": 256000,  
      "max_output_tokens": 8192,  
      "capabilities": {  
        "tool_use": true,  
        "embedding_use": true  
      }  
    }  
  ]  
}
```

```
        "long_context": true,
        "code": true
    },
},
{
    "id": "aurora-16b-assistant",
    "type": "text-completion",
    "context_window": 128000,
    "max_output_tokens": 4096,
    "capabilities": {
        "tool_use": true,
        "long_context": true,
        "code": true
    },
},
{
    "id": "aurora-3b-edge",
    "type": "text-completion",
    "context_window": 16000,
    "max_output_tokens": 2048,
    "capabilities": {
        "tool_use": false,
        "long_context": false,
        "code": true
    }
}
]
```

3. Chat Completions (Primary Endpoint)

Multi-turn conversation + tools + long context.

3.1 Create Chat Completion

POST /chat/completions

Request Body

```
{  
    "model": "aurora-72b-agent-v3",  
    "messages": [  
        {  
            "role": "system",  
            "content": "You are Aurora AI, a helpful technical assistant."  
        },  
        {  
            "role": "user",  
            "content": "Explain KV cache in transformers."  
        }  
    "tools": [  
        {  
            "name": "get_weather",  
            "description": "Fetch the weather for a given city.",  
            "input_schema": {  
                "type": "object",  
                "properties": {  
                    "location": { "type": "string" },  
                    "date": { "type": "string", "format": "date" }  
                },  
                "required": ["location"]  
            }  
        }  
    "tool_choice": "auto",  
    "max_tokens": 1024,  
    "temperature": 0.2,  
    "top_p": 0.9,  
    "n": 1,  
    "stream": false,  
    "metadata": {  
        "tenant_id": "acme-co",  
        "request_id": "abc-123"  
    }  
}
```

Parameters

- `model` (string, required)
 - e.g., `aurora-72b-agent-v3`

`messages` (array, required)

Array of message objects:

```
{  
  "role": "system" | "user" | "assistant" | "tool",  
  "name": "optional_tool_or_user_name",  
  "content": "text or tool result JSON-encoded as string"  
}
```

-
- `tools` (array, optional)
Tool definitions the model can call. See [§4](#) for full schema.
- `tool_choice` (string | object, optional)
 - `"auto"` – model decides when to call tools
 - `"none"` – tools disabled
 - `{"type": "tool", "name": "get_weather"}` – force a specific tool
- `max_tokens` (int, optional)
Max number of tokens in the completion.
- `temperature` (float, optional, default `0.7`)
- `top_p` (float, optional, default `1.0`)
- `n` (int, optional)
Number of completions to generate.
- `stream` (bool, optional)
Whether to stream tokens over SSE.

- **metadata** (object, optional)
Arbitrary key/value payload echoed back in logs.

Response

```
{  
  "id": "chatcmpl_01JABCDXYZ",  
  "object": "chat.completion",  
  "created": 1769097600,  
  "model": "aurora-72b-agent-v3",  
  "usage": {  
    "prompt_tokens": 134,  
    "completion_tokens": 256,  
    "total_tokens": 390  
  },  
  "choices": [  
    {  
      "index": 0,  
      "message": {  
        "role": "assistant",  
        "content": "Key-value (KV) cache is a memory structure used  
during transformer inference..."  
      },  
      "finish_reason": "stop",  
      "logprobs": null  
    }  
  ]  
}
```

Streaming Responses

If **stream: true**, Aurora returns an **SSE** stream:

```
POST /v1/chat/completions  
Accept: text/event-stream
```

Events:

```
data:  
{"id":"chatcmpl_...","object":"chat.completion.chunk","choices":[{"delta":{"role":"assistant","content":"KV "}, "index":0}]}  
data:  
{"id":"chatcmpl_...","object":"chat.completion.chunk","choices":[{"delta":{"content":"cache "}, "index":0}]}  
...  
data:  
{"id":"chatcmpl_...","object":"chat.completion.chunk","choices":[{"delta":{}, "finish_reason":"stop", "index":0}]}  
data: [DONE]
```

4. Tool Use & Agent Protocol

Aurora is trained to call tools using a structured schema similar to MCP/OpenAPI.

4.1 Tool Definition Format

Each tool object:

```
{  
  "name": "get_weather",  
  "description": "Fetches weather information for a city.",  
  "input_schema": {  
    "type": "object",  
    "properties": {  
      "location": {  
        "type": "string",  
        "description": "City name, e.g. 'Denver, CO'."  
      },  
      "date": {  
        "type": "string",  
        "format": "date",  
        "description": "Optional ISO 8601 date."  
      }  
    },  
    "required": ["location"]  
}
```

```
}
```

Constraints:

- `name` – must be a valid identifier: `[a-zA-Z0-9_-]+`
- `input_schema` follows JSON Schema (Draft 7-ish subset).

4.2 Tool Call Messages

When Aurora decides to call a tool, it responds with a message containing a `tool_calls` array.

Example assistant message:

```
{
  "role": "assistant",
  "content": null,
  "tool_calls": [
    {
      "id": "toolcall_01ABC",
      "type": "function",
      "name": "get_weather",
      "arguments": {
        "location": "Denver, CO",
        "date": "2025-11-22"
      }
    }
  ]
}
```

Notes:

- `arguments` is always a **JSON object**, already parsed for you.
- `content` is typically `null` for tool calls.

4.3 Tool Result Messages

Your application executes the tool and then **feeds the result back** as a new message:

```
{  
  "role": "tool",  
  "name": "get_weather",  
  "tool_call_id": "toolcall_01ABC",  
  "content": "{\"temperature_c\": 3.5, \"conditions\": \"snow\"}"  
}
```

Aurora then uses this tool result to continue the conversation.

4.4 Example End-to-End Tool Call

1. Client request:

```
{  
  "model": "aurora-72b-agent-v3",  
  "messages": [  
    { "role": "user", "content": "What's the weather in Denver today?" }  
  ],  
  "tools": [  
    {  
      "name": "get_weather",  
      "description": "Fetches today's weather.",  
      "input_schema": {  
        "type": "object",  
        "properties": {  
          "location": { "type": "string" }  
        },  
        "required": ["location"]  
      }  
    }  
  ],  
  "tool_choice": "auto"  
}
```

2. Aurora responds (tool call):

```
{  
  "choices": [  
    {  
      "message": {  
        "role": "assistant",  
        "tool_calls": [  
          {  
            "id": "toolcall_01ABC",  
            "type": "function",  
            "name": "get_weather",  
            "arguments": {  
              "location": "Denver, CO"  
            }  
          }  
        ]  
      }  
    }  
  ]  
}
```

3. Your app executes `get_weather` and sends another API call:

```
{  
  "model": "aurora-72b-agent-v3",  
  "messages": [  
    { "role": "user", "content": "What's the weather in Denver today?" },  
    {  
      "role": "assistant",  
      "tool_calls": [  
        {  
          "id": "toolcall_01ABC",  
          "type": "function",  
          "name": "get_weather",  
          "arguments": { "location": "Denver, CO" }  
        }  
      ]  
    },  
  ],  
  "function_call": "auto"
```

```
{  
    "role": "tool",  
    "name": "get_weather",  
    "tool_call_id": "toolcall_01ABC",  
    "content": "{\"temperature_c\": 3.5, \"conditions\": \"snow\"}"  
}  
]  
}
```

4. Aurora final answer:

```
{  
    "choices": [  
        {  
            "message": {  
                "role": "assistant",  
                "content": "In Denver today it's around 3.5°C with snow."  
            }  
        }  
    ]  
}
```

5. Plain Completions (Non-Chat)

For raw text-in, text-out tasks (no roles/messages).

5.1 Create Completion

POST /completions

Request

```
{  
    "model": "aurora-16b-assistant",  
    "prompt": "Write a function in Python that computes factorial of  
n.",  
    "max_tokens": 256,  
    "temperature": 0.3,
```

```
"top_p": 1.0,  
"n": 1,  
"stop": ["```"],  
"stream": false  
}
```

Response

```
{  
    "id": "cmpl_01ABA",  
    "object": "text.completion",  
    "created": 1769097680,  
    "model": "aurora-16b-assistant",  
    "usage": {  
        "prompt_tokens": 24,  
        "completion_tokens": 78,  
        "total_tokens": 102  
    },  
    "choices": [  
        {  
            "index": 0,  
            "text": "def factorial(n):\n                if n < 0:\n                    raise  
ValueError(\"n must be non-negative\")\n                result = 1\n                for i in  
range(2, n+1):\n                    result *= i\n                return result\n",  
            "finish_reason": "stop",  
            "logprobs": null  
        }  
    ]  
}
```

6. Embeddings

Generate vector embeddings for text (e.g., RAG, semantic search).

6.1 Create Embeddings

POST /embeddings

Request

```
{  
  "model": "aurora-embeddings-1024-v1",  
  "input": [  
    "Explain rotary position embeddings.",  
    "KV cache stores key and value tensors across layers."  
,  
  "encoding_format": "float"  
}
```

Response

```
{  
  "object": "list",  
  "model": "aurora-embeddings-1024-v1",  
  "data": [  
    {  
      "index": 0,  
      "embedding": [0.0231, -0.1043, 0.9932, ...]  
    },  
    {  
      "index": 1,  
      "embedding": [0.1189, 0.0021, 0.4023, ...]  
    }  
,  
  "usage": {  
    "prompt_tokens": 35,  
    "total_tokens": 35  
  }  
}
```

Parameters

- `model` – e.g. `aurora-embeddings-1024-v1`
- `input` – string or array of strings
- `encoding_format` – `"float"` or `"base64"`

7. Agents (High-Level Orchestrated Workflows)

A convenience layer that lets Aurora manage planning, tool calls, and final answers in a single endpoint.

7.1 Execute Agent

POST /agents/execute

Request

```
{  
  "agent_id": "default-aurora-agent",  
  "input": "Summarize this repo and generate a migration plan.",  
  "context_documents": [  
    {  
      "id": "doc_1",  
      "title": "README.md",  
      "content": "..."  
    },  
    {  
      "id": "doc_2",  
      "title": "MIGRATION_GUIDE.md",  
      "content": "..."  
    }  
  "max_steps": 8,  
  "return_plan": true,  
  "stream": false  
}
```

Response

```
{  
  "id": "agt_01XYZ",  
  "object": "agent.run",  
  "agent_id": "default-aurora-agent",  
  "created": 1769097800,  
  "steps": [  
    {  
      "id": "step_1",  
      "type": "planning",  
      "details": {  
        "plan": "Initial planning phase",  
        "status": "Planned"  
      }  
    },  
    {  
      "id": "step_2",  
      "type": "execution",  
      "details": {  
        "command": "git clone https://github.com/aurora-project/migration-guide.git",  
        "status": "Running"  
      }  
    },  
    {  
      "id": "step_3",  
      "type": "monitoring",  
      "details": {  
        "metrics": {"cpu_usage": 0.5, "memory_usage": 1.2},  
        "status": "Monitoring"  
      }  
    },  
    {  
      "id": "step_4",  
      "type": "finalization",  
      "details": {  
        "summary": "Migration guide generated successfully",  
        "status": "Completed"  
      }  
    }  
}
```

```
{  
    "index": 0,  
    "type": "plan",  
    "content": "<agent_plan>1) Read docs. 2) Extract current stack.  
3) Propose migration plan...</agent_plan>"  
},  
{  
    "index": 1,  
    "type": "tool_call",  
    "tool_name": "search_repo",  
    "arguments": { "query": "architecture overview" },  
    "result": { "matches": [ "doc_1", "doc_2" ] }  
}  
,  
]  
,  
"output": {  
    "role": "assistant",  
    "content": "Here's a summary of the repo and a suggested migration  
plan..."  
},  
"usage": {  
    "prompt_tokens": 300,  
    "completion_tokens": 700,  
    "total_tokens": 1000  
}  
}
```

8. Rate Limiting & Quotas

8.1 Headers

Every response includes rate-limit headers:

X-Aurora-Limit-Requests: 3000
X-Aurora-Remaining-Requests: 2789
X-Aurora-Reset-Requests: 1769101200

X-Aurora-Limit-Tokens: 300000

X-Aurora-Remaining-Tokens: 284523

X-Aurora-Reset-Tokens: 1769101200

- *Limit-* – quota per rolling window
- *Remaining-* – remaining allowance
- *Reset-* – UNIX timestamp when quota resets

8.2 Rate-Limit Error

```
{  
  "error": {  
    "code": "rate_limit_exceeded",  
    "message": "Token quota exceeded for this window.",  
    "details": {  
      "retry_after_seconds": 30  
    }  
  }  
}
```

9. Errors

All error responses use a common envelope:

```
{  
  "error": {  
    "code": "invalid_request",  
    "message": "The 'model' field is required.",  
    "details": {  
      "field": "model",  
      "location": "body"  
    }  
  }  
}
```

9.1 Common Error Codes

- `invalid_request` – malformed JSON, missing fields, invalid parameter range
 - `unauthorized` – invalid API key
 - `forbidden` – insufficient permissions
 - `not_found` – unknown model or resource
 - `rate_limit_exceeded` – quotas exceeded
 - `server_error` – internal Aurora error
 - `overloaded` – temporary capacity issue, retry with backoff
-

10. Safety Controls

10.1 Safe-Mode Parameters (Chat & Completions)

You can control safety behavior:

```
{  
  "safety": {  
    "level": "standard",  
    "allow_deep_technical": true,  
    "blocked_categories": ["self_harm", "weapons"]  
  }  
}
```

- `level` – "strict" | "standard" | "relaxed" (tenant-configurable defaults)
- `blocked_categories` – override to explicitly block certain content domains

10.2 Safety Metadata in Responses

```
{
```

```
"choices": [
  {
    "message": { "role": "assistant", "content": "..." },
    "safety": {
      "policy_applied": true,
      "categories": {
        "self_harm": "none",
        "harassment": "none",
        "hate": "none"
      }
    }
  }
]
```