

Aurora AI — Token / Parameter Specifications & Vector Mechanics (Plain-Text Version)

Below is a fully text-safe, technical, researcher-grade breakdown.

1. Tokenization Specification

1.1 Text Normalization

Aurora's tokenizer applies the following normalization steps:

1. Unicode normalization: NFC
2. Case-preserving (no forced lowercase)
3. Whitespace handling: collapse multiple spaces, preserve newlines and tabs
4. Control characters removed except newline and tab
5. Full-width → half-width conversion for CJK text

1.2 Subword Model

Aurora uses a BPE or Unigram LM tokenizer.

- Vocabulary size: typically between 64k and 256k
- Joint text + code training
- Optimized for identifiers, URLs, snake_case, camelCase, JSON keys

Pre-tokenization

Splitting happens on:

- Spaces
- Standard punctuation
- Unicode categories Zs and P*
- Punctuation retained as standalone subword candidates

1.3 Vocabulary Layout (Plain Text IDs)

Recommended layout:

- 0: <pad>
- 1: <bos>
- 2: <eos>
- 3: <unk>
- 4: <mask>
- 5–31: reserved system/control tokens
- 32–V–1: learned subword tokens

Example special token subset: {0..31}

1.4 Sequence Length & Packing

- Maximum context length: L_max (example: 16,384 tokens)
 - Input tensor shape: [batch_size, L_max]
 - Attention mask shape: [batch_size, L_max, L_max]
 - Sequences may be "packed" using <sep> and block-diagonal attention masks
-

2. Parameter Specification

This section uses plain-text arithmetic instead of formulas.

2.1 Definitions

Let:

- V = vocabulary size
- d_{model} = transformer hidden dimension
- n_{layer} = number of transformer blocks
- n_{head} = number of attention heads
- d_{head} = dimension per head
- d_{ff} = hidden dimension inside the feedforward network

Relationship:

- $d_{\text{model}} = n_{\text{head}} * d_{\text{head}}$
-

2.2 Parameter Counts (Plain-Text Descriptions)

2.2.1 Embedding Matrix

Shape: $[V, d_{\text{model}}]$

Parameter count: $V * d_{\text{model}}$

2.2.2 Positional or Rotary Embeddings

- RoPE has negligible learned parameters
- Learned positional embeddings (optional) would have shape: `[L_max, d_model]`

2.2.3 Transformer Block (per layer)

Each block includes:

- LayerNorm
- Multi-head self-attention
- MLP / feedforward network

Attention Parameters Per Layer

Each block contains 4 projection matrices:

- `W_Q`: `[d_model, d_model]`
- `W_K`: `[d_model, d_model]`
- `W_V`: `[d_model, d_model]`
- `W_O`: `[d_model, d_model]`

Total attention parameters per layer:

- `4 * (d_model * d_model)`

MLP / FFN Parameters Per Layer

Standard 2-layer FFN:

- `W1`: `[d_model, d_ff]`
- `W2`: `[d_ff, d_model]`
- `Total = 2 * d_model * d_ff`

Gated FFN (e.g., SwiGLU):

- $\text{Total} = 4 * \text{d_model} * \text{d_ff}$

LayerNorm Parameters

Very small:

- Each LayerNorm has $2 * \text{d_model}$ parameters
- Two LayerNorms per block → approx $4 * \text{d_model}$ per block

Usually considered negligible at this scale.

2.2.4 Output Projection / LM Head

If **tied** with embeddings:

- No additional parameters.

If untied:

- Same shape as embedding matrix: $[\text{V}, \text{ d_model}]$

2.2.5 Total Parameter Count

Approximate:

- $\text{Total} \approx (\text{V} * \text{d_model}) + \text{n_layer} * (4 * \text{d_model}^2 + 4 * \text{d_model} * \text{d_ff})$

(Assuming gated MLP and tied embeddings.)

2.3 Example: Aurora-32B Configuration (Plain Text Version)

Example values:

- $V = 128,000$
- $d_{model} = 6,144$
- $n_{layer} = 48$
- $n_{head} = 48$
- $d_{head} = 128$
- $d_{ff} = 24,576$ (which is $4 * d_{model}$)
- $L_{max} = 16,384$
- $\text{MLP} = \text{gated (SwiGLU)}$

Parameter estimation:

- Embeddings:

$$128,000 * 6,144 \approx 787 \text{ million}$$

- Per-layer attention:

$$4 * (6,144 * 6,144) \approx 151 \text{ million}$$

- Per-layer gated FFN:

$$4 * (6,144 * 24,576) \approx 604 \text{ million}$$

- Total per layer:

$$\text{approx } 755 \text{ million}$$

- All 48 layers:

$$48 * 755M \approx 36B \text{ parameters}$$

The resulting design is marketed as ~32B due to pruning, weight sharing, and quantization optimizations.

3. How Vectors Work in Aurora AI (Plain-Text Version)

This removes all symbolic notation and uses direct prose and pseudo-code.

3.1 Token → Embedding Vectors

Given input token IDs:

- Input shape: [batch_size, sequence_length]

Step 1: Look up embeddings:

- $H[\text{batch}, \text{position}] = \text{embedding_matrix}[\text{token_id}]$
- Output shape: [batch_size, sequence_length, d_model]

Step 2: Apply RoPE inside the attention mechanism, not on embeddings directly.

3.2 Self-Attention (Plain Text Explanation)

For each transformer layer:

1. Project embeddings into Query, Key, Value matrices:

- $Q = H * W_Q$
- $K = H * W_K$
- $V = H * W_V$

Each is reshaped to:

- [sequence_length, n_head, d_head]

2. Apply RoPE rotations to Q and K.
3. Compute attention scores:
 - `score[t][u][h] = dot(Q[t][h], K[u][h]) / sqrt(d_head)`

Causal rule:

- Positions after t are masked (ignored).
4. Convert scores to attention weights using softmax.
 5. Compute attention output:
 - `O[t][h] = sum_over_u (attention_weight[t][u][h] * V[u][h])`
 6. Concatenate all heads and project through W_O.
 7. Add residual connection to form intermediate output.

3.3 MLP / FFN (Plain Text)

Each token's vector goes through:

1. LayerNorm
2. Two linear layers (or three for gated MLP)
3. Activation (Swish for SwiGLU)
4. Residual connection

Equivalent pseudo-code:

- `Z = LayerNorm(H)`

- $U = Z * W1_u$
 - $V = Z * W1_v$
 - $Gated = \text{Swish}(U)$ elementwise-multiplied by V
 - $M = Gated * W2$
 - $H_{\text{out}} = H + M$
-

3.4 Final Logits (Plain Text)

1. Apply final LayerNorm.
 2. Multiply by the transpose of the embedding matrix.
 3. Apply softmax to get probabilities for next token.
-

4. External Vector Interfaces (Retrieval, Embeddings, RAG)

4.1 Token Embeddings

Use hidden states from last layer:

- `shape: [batch, sequence_length, d_model]`

4.2 Sequence Embeddings

Common pooling options:

- Take the hidden state at position of <bos> token
- Mean pooling: average all token vectors

- Weighted pooling

Example mean pooling:

- `embedding = average(H[last_layer][tokens])`

4.3 Embedding Projection Head

Optional:

- `projected_embedding = normalize(embedding * W_emb)`

Helps for RAG vector stores.

4.4 Similarity Measurement

Cosine similarity:

- `cosine = dot(a, b) / (norm(a) * norm(b))`

5. Vector Steering / Directional Arithmetic

Attributes can be represented as direction vectors by averaging embeddings of examples.

Example:

- `mean_A = average(embeddings from set A)`
- `mean_B = average(embeddings from set B)`
- `direction = mean_A - mean_B`
- `new_vector = z + lambda * direction`

Used experimentally to explore geometry of representation space.

6. Numeric Precision

Training:

- Weights: bf16 or fp16
- Activations: bf16, fp16, or fp8
- Accumulators: fp32

Serving:

- Weights: int8, nf4, or fp16

Embeddings returned in fp32 or fp16

-